

NuFast: Efficient Three-Flavor Neutrino Oscillation Probabilities in Rust

Baalateja Kataru

Planckeon Labs
baalateja@planckeon.org

February 2026

Abstract. We present `nufast`, a Rust implementation of the NuFast algorithm for computing three-flavor neutrino oscillation probabilities in vacuum and matter. Our implementation achieves performance competitive with optimized C++ code, with 61 ns for vacuum and 95 ns for matter calculations per energy point—approximately 27% faster than C++ for matter effects. The crate is published on crates.io and includes WebAssembly bindings for browser-based applications. This work enables high-performance neutrino physics calculations in modern software ecosystems.

1 Introduction

Neutrino oscillation is a quantum mechanical phenomenon where neutrinos change flavor as they propagate through space. Accurate and efficient computation of oscillation probabilities is essential for analyzing data from neutrino experiments such as DUNE, Hyper-Kamiokande, and JUNO.

The NuFast algorithm, developed by Denton and Parke [1], provides a computationally efficient method for calculating three-flavor oscillation probabilities including matter effects. The algorithm uses the Eigenvalue- Eigenvector Identity (EEI) which only requires diagonalizing 2×2 matrices (solving quadratics rather than cubics for eigenvectors). An additional optimization is that the square root for the quadratic need not be computed to obtain probabilities. The algorithm also uses an initial eigenvalue approximation that propagates to other eigenvalues in an optimal order, with iterative Newton-Raphson refinement for arbitrary precision.

NuFast has been adopted by major experimental collaborations: it is implemented in the MaCH3 framework (the primary reweighting code for T2K and other US/Japan experiments) and JUNO analysis pipelines, achieving “dramatic speed ups—close to an order of magnitude—over other ‘optimized’ algorithms” [1].

We present `nufast`, a Rust port of this algorithm, designed to leverage Rust’s memory safety guarantees and zero-cost abstractions while achieving performance comparable to or better than the original implementations.

1.1 Motivation

This project represents the culmination of several years of work on neutrino oscillation phenomenology.

1.1.1 Undergraduate Research (2022–2023)

The author's undergraduate capstone thesis at Krea University, supervised by Dr. Sushant Raut, explored the “Interplay between Neutrino Oscillations and Linear Algebra.” The research investigated applications of the Eigenvalue- Eigenvector Identity (EEI, also called the Rosetta identity) discovered by Denton, Parke, and Zhang [2], as well as the Adjugate Identity [3], to streamline symbolic calculations of oscillation probabilities.

The goal was to derive novel series expansions of oscillation probabilities in matter up to second order in the mass hierarchy parameter α only—as opposed to second order in both α and s_{13} as in the influential work of Akhmedov et al. [4]. Using Mathematica and the Cayley-Hamilton formalism, the author explored whether these linear algebra identities could simplify the analytic calculation. While the symbolic expansions proved too computationally intensive (repeatedly crashing Mathematica), the numerical implementations were successful, resulting in `pytrino`—a Python/Cython library published on PyPI that computes oscillation probabilities using the EEI and Adjugate identities.

1.1.2 Postgraduate Research (2023–2024)

The author continued this work during a postgraduate program at Krea University, pivoting to investigate neutrino oscillations on quantum computers using Hamiltonian simulation (Trotter-Suzuki decomposition) and quantum machine learning approaches. This work successfully reproduced published results [5, 6] using IBM’s Qiskit framework.

1.1.3 Correspondence with Dr. Denton (October 2024)

In October 2024, the author reached out to Dr. Peter Denton regarding research opportunities in neutrino physics, describing the prior work on the EEI and the challenges encountered with 3+1 sterile neutrino models. Dr. Denton explained that while the EEI is powerful for three-flavor oscillations (requiring only 2×2 matrix diagonalization), extending to four flavors requires solving cubic equations for eigenvectors—“analytically much much worse, and also numerically somewhat unstable.”

Dr. Denton recommended the NuFast algorithm: “We recently put together what we think is the optimal algorithm” (arXiv:2405.02400), noting its adoption in the MaCH3 framework and JUNO pipelines. He reported 100 ns per probability calculation on his laptop.

1.1.4 Rust Implementation (2024–2026)

The decision to port NuFast to Rust was motivated by:

1. **Performance exploration:** Testing whether Rust’s ownership model and LLVM backend could match or exceed C++. Our implementation achieves 95 ns for matter calculations—beating Dr. Denton’s quoted benchmark.
2. **Modern tooling:** Enabling WebAssembly deployment for browser applications
3. **Educational visualization:** Creating “Imagining the Neutrino,” an interactive web-based tool for teaching neutrino oscillation concepts

2 Algorithm Overview

The NuFast algorithm computes the full 3×3 oscillation probability matrix $P_{\alpha\beta}$ for neutrino flavor transitions $\nu_\alpha \rightarrow \nu_\beta$ where $\alpha, \beta \in \{e, \mu, \tau\}$.

2.1 Vacuum Oscillations

In vacuum, the oscillation probability depends on:

- Mixing angles: $\theta_{12}, \theta_{13}, \theta_{23}$
- CP-violating phase: δ
- Mass-squared differences: $\Delta m_{21}^2, \Delta m_{31}^2$
- Baseline L and energy E

The algorithm computes exact probabilities using trigonometric identities without matrix diagonalization.

2.2 Matter Effects

For propagation through matter with constant density ρ and electron fraction Y_e , the Mikheyev-Smirnov-Wolfenstein (MSW) effect modifies the effective mixing parameters. NuFast uses:

1. An initial estimate from the DMP (Denton-Minakata-Parke) approximation
2. Optional Newton-Raphson iterations to improve precision

The number of Newton iterations N_{Newton} controls the trade-off between speed and precision. For most long-baseline experiments, $N_{\text{Newton}} = 0$ or 1 is sufficient.

3 Implementation

Our Rust implementation provides a simple API:

```
use nufast::{VacuumParameters, MatterParameters};
use nufast::{probability_vacuum_lbl, probability_matter_lbl};

// Vacuum oscillation
let params = VacuumParameters::nufit52_no(1300.0, 2.5);
let probs = probability_vacuum_lbl(&params);
println!("P(mu-e) = {}", probs.Pme);

// Matter oscillation
let params = MatterParameters::nufit52_no(1300.0, 2.5, 3.0, 0.5, 0);
let probs = probability_matter_lbl(&params);
```

The crate is published on crates.io at <https://crates.io/crates/nufast>.

3.1 WebAssembly Support

For web applications, we provide `nufast-wasm`, which compiles the core physics engine to WebAssembly. This enables browser-based neutrino physics tools with near-native performance. The compiled WASM module is approximately 32 KB gzipped.

4 Benchmark Methodology

All benchmarks were performed on a system with AMD Ryzen CPU running WSL2 on Windows. Each benchmark:

1. Iterates over 10 million (10^7) calculations
2. Varies energy from 0.5–5.0 GeV (DUNE-like range)
3. Uses standard oscillation parameters (NuFIT 5.2)
4. Includes a sink variable to prevent compiler optimization

The Rust benchmarks used the Criterion library with statistical analysis. C++, Fortran, and Python benchmarks used high-resolution timing.

5 Results

Language	Vacuum	N=0	N=1	N=2	N=3
Rust	61 ns	95 ns	106 ns	113 ns	117 ns
C++	49 ns	130 ns	143 ns	154 ns	164 ns
Fortran	51 ns	107 ns	123 ns	146 ns	167 ns
Python	14,700 ns	21,900 ns	21,200 ns	18,500 ns	16,300 ns

Table 1: Single-point oscillation probability computation times (ns per call). N refers to the number of Newton-Raphson iterations.

5.1 Key Findings

Comparison	Vacuum	Matter (N=0)	Notes
Rust vs C++	+24%	-27%	Rust faster for matter
Rust vs Fortran	+20%	-11%	Rust faster for matter
Rust vs Python	$\times 241$	$\times 230$	Compiled vs interpreted

Table 2: Relative performance (negative = Rust is faster). Rust shows a significant advantage for matter calculations.

5.1.1 Rust is Faster for Matter Calculations

The most surprising result is that Rust outperforms C++ by **27%** for matter oscillations. This is likely due to:

1. **Better loop optimization:** LLVM's optimization of the Newton iteration
2. **Stricter aliasing rules:** Rust's ownership model enables more aggressive optimization
3. **Modern code patterns:** The Rust implementation uses idiomatic patterns that optimize well

5.1.2 Vacuum Performance

For vacuum calculations, C++ and Fortran are 20% faster than Rust. This is a smaller and simpler computation where traditional numerical languages have an edge.

5.1.3 Python Performance

As expected, Python is approximately 240× slower than compiled languages. However, for interactive exploration or small-scale calculations, this overhead is acceptable.

5.2 Throughput

Language	Vacuum	Matter (N=0)
Rust	17.5 M/s	10.5 M/s
C++	20.3 M/s	7.7 M/s
Fortran	19.7 M/s	9.4 M/s
Python	0.07 M/s	0.05 M/s

Table 3: Throughput in millions of probability calculations per second.

6 Applications

6.1 Interactive Visualization

The `nufast-wasm` module powers “Imagining the Neutrino,” an interactive web-based visualization of neutrino oscillations: <https://planckeon.github.io/itn/>

With WASM, the visualization computes 400-point energy spectra in real-time as users adjust parameters.

6.2 Integration with Analysis Frameworks

The crate can be integrated with Rust-based physics analysis pipelines, or called from Python via PyO3 bindings (future work).

7 Conclusion

We have demonstrated that Rust provides a viable and performant platform for computational neutrino physics. The `nufast` crate achieves:

- **27% speedup** over C++ for matter calculations
- **Competitive performance** for vacuum calculations
- **Memory safety** guarantees without runtime overhead
- **WebAssembly support** for browser applications
- **Published on crates.io** for easy integration

The combination of performance, safety, and modern tooling makes Rust an attractive choice for future neutrino physics software development.

Acknowledgments

The author thanks Dr. Peter B. Denton (Brookhaven National Laboratory) for recommending the NuFast algorithm and providing guidance on optimal approaches to neutrino oscillation calculations. This work builds on the NuFast algorithm developed by Denton and Parke, with original implementations available at <https://github.com/PeterDenton/NuFast>.

The author also acknowledges the neutrino physics education received during undergraduate research at Krea University under Dr. Sushant Raut, which provided the theoretical foundation for this implementation.

Code Availability

The `nufast` crate is open source under the MIT license:

- Crates.io: <https://crates.io/crates/nufast>
- GitHub: <https://github.com/planckeon/nufast>
- Documentation: <https://docs.rs/nufast>

All benchmark implementations (Rust, C++, Fortran, Python) are included in the repository under `benchmarks/`.

References

1. P. B. Denton, “Neutrino Oscillation Probabilities: A Compact Multi-algorithm Approach,” arXiv:2405.02400 (2024).
2. P. B. Denton, S. J. Parke, and X. Zhang, “Eigenvector-based approach to neutrino oscillation probabilities in matter,” Phys. Rev. D 101, 093001 (2020). arXiv:1907.02534
3. A. Abdullahi and S. J. Parke, “Eigenvalue-independent eigenvectors and the adjugate of a matrix,” arXiv:2212.12565 (2022).
4. E. K. Akhmedov, R. Johansson, M. Lindner, T. Ohlsson, and T. Schwetz, “Series expansions for three-flavor neutrino oscillation probabilities in matter,” JHEP 04, 078 (2004). arXiv:hep-ph/0402175
5. C. A. Argüelles and B. J. P. Jones, “Neutrino Oscillations in a Quantum Computer,” Phys. Rev. D 99, 096005 (2019). arXiv:1904.10559
6. J. Turro et al., “A Quantum simulation of neutrino-matter collective oscillations,” arXiv:2111.05401 (2021).
7. DUNE Collaboration, “Long-Baseline Neutrino Facility (LBNF) and Deep Underground Neutrino Experiment (DUNE) Conceptual Design Report,” arXiv:1601.05471
8. Hyper-Kamiokande Proto-Collaboration, “Hyper-Kamiokande Design Report,” arXiv:1805.04163
9. JUNO Collaboration, “Neutrino Physics with JUNO,” J. Phys. G 43, 030401 (2016). arXiv:1507.05613