

NuFast: Efficient Three-Flavor Neutrino Oscillation Probabilities in Rust

Baalateja Kataru

Planckeon Labs
baalateja@planckeon.org

February 3, 2026

Abstract

We present `nufast`, a Rust implementation of the NuFast algorithm for computing three-flavor neutrino oscillation probabilities in vacuum and constant-density matter. Our implementation achieves performance competitive with optimized C++ code: ~ 61 ns for vacuum and ~ 95 ns for matter calculations per energy point—approximately **27% faster** than C++ for matter effects. This performance advantage stems from LLVM’s aggressive optimization of Rust’s ownership-based memory model during Newton-Raphson iterations. The crate is published on crates.io and includes WebAssembly bindings enabling browser-based applications. We also provide `VacuumBatch`, an optimized API for batch calculations that pre-computes mixing matrix elements, achieving 45% speedup for energy spectrum computations. This work enables high-performance neutrino physics calculations in modern software ecosystems while maintaining memory safety guarantees.

Keywords: neutrino oscillations, matter effects, MSW effect, NuFast algorithm, Rust, WebAssembly, high-performance computing

1. Introduction

Neutrino oscillation is a quantum mechanical phenomenon where neutrinos change flavor as they propagate through space. The discovery of neutrino oscillations—implying nonzero neutrino masses—represents physics beyond the Standard Model and was recognized with the 2015 Nobel Prize in Physics. Accurate and efficient computation of oscillation probabilities is essential for analyzing data from current and next-generation experiments including DUNE, Hyper-Kamiokande, and JUNO.

The transition probability from flavor α to flavor β after propagating distance L with energy E is given by:

$$P_{\alpha\beta} = |\langle \nu_\beta | \nu_\alpha(L) \rangle|^2 = \sum_{i,j} U_{\alpha i}^* U_{\beta i} U_{\alpha j} U_{\beta j}^* e^{-i\Delta m_{ij}^2 L / 2E} \quad (1)$$

where U is the Pontecorvo-Maki-Nakagawa-Sakata (PMNS) mixing matrix and $\Delta m_{ij}^2 = m_i^2 - m_j^2$ are mass-squared differences.

In matter, the Mikheyev-Smirnov-Wolfenstein (MSW) effect modifies the effective mixing parameters through coherent forward scattering of electron neutrinos on electrons. Computing these modified probabilities efficiently has been a longstanding challenge in neutrino phenomenology.

1.1. The NuFast Algorithm

The NuFast algorithm, developed by Denton and Parke [1], provides a computationally optimal method for three-flavor oscillation probabilities. Key innovations include:

1. **Eigenvalue-Eigenvector Identity (EEI):** Avoids cubic equation solving by using only 2×2 matrix diagonalization (quadratic equations)
2. **Square root elimination:** The quadratic's discriminant square root cancels in probability expressions
3. **Optimal eigenvalue ordering:** Initial DMP approximation propagates to other eigenvalues efficiently
4. **Newton-Raphson refinement:** Optional iterations for arbitrary precision

Remark. NuFast has been adopted by major collaborations: it is implemented in MaCH3 (the primary reweighting framework for T2K and other US/Japan experiments) and JUNO analysis pipelines, achieving “dramatic speed ups—close to an order of magnitude—over other ‘optimized’ algorithms” [1].

1.2. Motivation and Historical Context

This implementation represents the culmination of several years of work on neutrino oscillation phenomenology.

1.2.1. Undergraduate Research (2022–2023)

The author’s undergraduate capstone thesis at Krea University, supervised by Dr. Sushant Raut, explored the *Interplay between Neutrino Oscillations and Linear Algebra*. The research investigated applications of the Eigenvalue-Eigenvector Identity (also called the Rosetta identity) [2] and the Adjugate Identity [3] to streamline symbolic calculations of oscillation probabilities.

The goal was to derive novel series expansions of oscillation probabilities in matter up to second order in the mass hierarchy parameter $\alpha \equiv \Delta m_{21}^2 / \Delta m_{31}^2$ only—as opposed to second order in both α and $\sin \theta_{13}$ as in Akhmedov et al. [4]. Using Mathematica and the Cayley-Hamilton formalism, the author explored whether these linear algebra identities could simplify the analytic calculation. While the symbolic expansions proved computationally intractable (repeatedly exhausting available memory), numerical implementations were successful, resulting in `pytrino`—a Python/Cython library published on PyPI.

1.2.2. Postgraduate Research (2023–2024)

The author continued this work during a postgraduate program, pivoting to investigate neutrino oscillations on quantum computers using Hamiltonian simulation (Trotter-Suzuki decomposition) and quantum machine learning approaches. This work reproduced published results [5], [6] using IBM’s Qiskit framework.

1.2.3. Correspondence with Dr. Denton (October 2024)

In October 2024, the author consulted Dr. Peter Denton regarding research directions in neutrino physics, describing prior work on the EEI and challenges with 3+1 sterile neutrino extensions. Dr. Denton explained that while the EEI is powerful for three flavors (2×2 diagonalization), four-flavor oscillations require cubic eigenvector equations—“analytically much much worse, and also numerically somewhat unstable.”

Dr. Denton recommended the NuFast algorithm, noting approximately 100 ns per probability calculation on his laptop. This benchmark became a target for our Rust implementation.

1.2.4. Precursor Rust Implementations (2024)

Before implementing the full NuFast algorithm, the author developed several prototype implementations:

- **rustrino** (August 2024): Initial exploration of Rust for neutrino physics, establishing the development environment and build toolchain.
- **nosc** (October 2024): A working two-flavor oscillation engine implementing the standard vacuum and matter formulas. This served as a testbed for understanding Rust’s numerical performance characteristics and API design patterns for physics libraries.

These incremental steps informed the design decisions in the final `nufast` implementation, particularly regarding:

- Parameter struct design (`VacuumParameters`, `MatterParameters`)
- Result struct layout for probability triplets
- Benchmark methodology using Criterion

The complete chronology of the author's neutrino software development is:

Date	Project	Description
May 2023	<code>pytrino</code>	Python/Cython library (undergrad thesis)
Aug 2024	<code>rustrino</code>	First Rust prototype
Sep 2024	<code>nufast</code>	NuFast algorithm port (Rust)
Oct 2024	<code>nosc</code>	Two-flavor Rust implementation
Feb 2026	<code>nufast-zig</code>	Zig port with SIMD support
Feb 2026	<code>nufast-wasm</code>	WebAssembly + TypeScript bindings

Table 1: Chronology of the author's neutrino oscillation software.

2. Algorithm Details

2.1. Vacuum Oscillations

In vacuum, the oscillation probability depends on:

- **Mixing angles:** $\theta_{12}, \theta_{13}, \theta_{23}$
- **CP-violating phase:** δ_{CP}
- **Mass-squared differences:** $\Delta m_{21}^2, \Delta m_{31}^2$
- **Baseline and energy:** L (km), E (GeV)

Definition (PMNS Matrix). The PMNS matrix in the standard parameterization is:

$$U = \begin{pmatrix} c_{12}c_{13} & s_{12}c_{13} & s_{13}e^{-i\delta} \\ -s_{12}c_{23} - c_{12}s_{23}s_{13}e^{i\delta} & c_{12}c_{23} - s_{12}s_{23}s_{13}e^{i\delta} & s_{23}c_{13} \\ s_{12}s_{23} - c_{12}c_{23}s_{13}e^{i\delta} & -c_{12}s_{23} - s_{12}c_{23}s_{13}e^{i\delta} & c_{23}c_{13} \end{pmatrix} \quad (2)$$

where $c_{ij} = \cos \theta_{ij}$ and $s_{ij} = \sin \theta_{ij}$.

The vacuum algorithm computes exact probabilities using trigonometric identities without matrix exponentiation or diagonalization.

2.2. Matter Effects

For propagation through matter with constant electron density N_e , the effective Hamiltonian acquires a matter potential:

$$H = H_{\text{vacuum}} + \text{diag}(a, 0, 0), \quad a = \sqrt{2}G_F N_e = 7.63 \times 10^{-5}(\rho Y_e)[\text{eV}^2 / \text{GeV}] \quad (3)$$

where ρ is the matter density in g/cm³ and Y_e is the electron fraction.

NuFast uses:

1. **DMP approximation:** Initial eigenvalue estimate from Denton-Minakata-Parke [7]

2. Newton-Raphson refinement: N_{Newton} iterations for improved precision

Remark. For long-baseline experiments like DUNE (1300 km, 2.5 GeV, $\rho \approx 2.8 \text{ g/cm}^3$), $N_{\text{Newton}} = 0$ provides sub-percent accuracy. Higher precision is available with $N_{\text{Newton}} = 1$ or 2.

3. Implementation

3.1. Core API

Our Rust implementation provides an ergonomic API:

```
use nufast::{VacuumParameters, MatterParameters};
use nufast::{probability_vacuum_lbl, probability_matter_lbl};

// Vacuum oscillation with NuFIT 5.2 parameters
let params = VacuumParameters::nufit52_no(1300.0, 2.5);
let probs = probability_vacuum_lbl(&params);
println!("P(vμ → ve) = {:.4}", probs.Pme);

// Matter oscillation
let params = MatterParameters::nufit52_no(
    1300.0, // L (km)
    2.5, // E (GeV)
    2.8, // ρ (g/cm³)
    0.5, // Ye
    0 // N_Newton
);
let probs = probability_matter_lbl(&params);
```

3.2. VacuumBatch Optimization

For batch calculations (e.g., computing energy spectra), we provide `VacuumBatch` which pre-computes mixing matrix elements:

```
use nufast::VacuumBatch;

let batch = VacuumBatch::nufit52_no();
let spectrum = batch.spectrum(1300.0, 0.5, 5.0, 1000);
// 1000-point spectrum in ~72 μs
```

Key Result. `VacuumBatch` achieves **45% speedup** over repeated single-point calls by pre-computing all nine $|U_{\alpha i}|^2$ elements and the Jarlskog invariant once, then reusing them across all energy/baseline points.

3.3. WebAssembly Support

The `nufast-wasm` crate compiles the physics engine to WebAssembly:

```
import init, { wasmCalculateEnergySpectrum } from 'nufast-wasm';

await init();
const spectrum = wasmCalculateEnergySpectrum({
    theta12_deg: 33.44,
    theta13_deg: 8.57,
    // ... other parameters
}, 1300, 0.5, 5.0, 200);
```

The compiled WASM module is approximately **32 KB gzipped**, enabling browser-based neutrino physics with near-native performance.

3.4. Zig Implementation

We also provide a Zig implementation that achieves the highest performance:

```
const nufast = @import("nufast");

// Vacuum oscillation
const params = nufast.VacuumParams.default;
const probs = nufast.vacuumProbability(params, 1300.0, 2.5);
// probs[1][0] = P(vμ → ve)

// Batch calculations with pre-computed matrix
const batch = nufast.VacuumBatch.init(params);
for (energies) |E| {
    const p = batch.probabilityAt(1300.0, E);
}

// SIMD: 4 energies simultaneously (f64)
var E_vec: nufast.F64Vec = .{ 1.0, 2.0, 3.0, 4.0 };
const p_vec = nufast.vacuumProbabilitySimd(batch, L, E_vec);

// f32 SIMD: 8 energies simultaneously (2x throughput)
const batch_f32 = nufast.VacuumBatchF32.fromF64(batch);
var E_f32: nufast.F32Vec = .{ 1, 1.5, 2, 2.5, 3, 3.5, 4, 4.5 };
const p_f32 = nufast.vacuumProbabilitySimdF32(batch_f32, L, E_f32);
```

The Zig implementation includes:

- **Native SIMD** via `@Vector` types (4×f64 or 8×f32 on AVX2)
- **Zero heap allocations** in hot paths
- **f32 mode** for 2× SIMD lanes when full precision isn't needed
- **MatterBatch** for pre-computed constant-density calculations
- **Anti-neutrino mode** with sign-flipped matter potential and δCP
- **Experiment presets** for DUNE, T2K, NOvA, Hyper-K, and JUNO

3.4.1. Experiment Presets

The Zig implementation includes pre-configured parameters for common experiments:

```
const nufast = @import("nufast");

// Use a preset directly
const dune = nufast.experiments.dune;
const probs = nufast.matterProbability(dune.toMatterParams(), dune.L, dune.E);

// Available: t2k, nova, dune, hyper_k, juno
```

Each preset provides baseline L , peak energy E , and appropriate matter density for the experiment's path through Earth.

3.5. Anti-Neutrino Support

Both Rust and Zig implementations support anti-neutrino calculations:

```
// Rust: anti-neutrino mode
let mut params = MatterParameters::nufit52_no(1300.0, 2.5);
params.antineutrino = true;
let probs = probability_matter_lbl(&params);

// Zig: anti-neutrino mode
var params = nufast.MatterParams.default;
params.antineutrino = true;
const probs = nufast.matterProbability(params, 1300.0, 2.5);
```

For anti-neutrinos, two sign flips are applied:

1. **CP phase:** $\delta \rightarrow -\delta$ (CPT conjugation)
2. **Matter potential:** $A \rightarrow -A$ (opposite sign for anti-particles)

This enables direct calculation of CP asymmetries: $A_{\text{CP}} = P(\nu_\mu \rightarrow \nu_e) - P(\bar{\nu}_\mu \rightarrow \bar{\nu}_e)$

4. Benchmark Methodology

All benchmarks were performed on AMD Ryzen (WSL2/Windows). Methodology:

- **Iterations:** 10 million (10^7) calculations per measurement
- **Energy range:** 0.5–5.0 GeV (DUNE-like parameters)
- **Parameters:** NuFIT 5.2 best-fit values [8]
- **Anti-optimization:** Sink variables prevent dead code elimination

Rust benchmarks use Criterion with statistical analysis. C++, Fortran, and Python use high-resolution timing with standard deviation over 10 runs.

5. Results

Language	Vacuum	N=0	N=1	N=2	N=3
Zig	25.6 ns	73.3 ns	77.7 ns	81.1 ns	82.4 ns
Rust	61 ns	95 ns	106 ns	113 ns	117 ns
C++	49 ns	130 ns	143 ns	154 ns	164 ns
Fortran	51 ns	107 ns	123 ns	146 ns	167 ns
Python	14,700 ns	21,900 ns	21,200 ns	18,500 ns	16,300 ns

Table 2: Single-point oscillation probability timing (ns/call). N = Newton-Raphson iterations. Bold indicates fastest per column.

Comparison	Vacuum	Matter (N=0)	Interpretation
Zig vs Rust	-58%	-23%	Zig is fastest
Zig vs C++	-48%	-44%	Zig dominates
Rust vs C++	+24%	-27%	Rust faster for matter
Rust vs Fortran	+20%	-11%	Rust competitive
Zig vs Python	×574	×299	Compiled advantage

Table 3: Relative performance (negative = faster).

5.1. Key Findings

5.1.1. Zig Achieves Fastest Performance

The Zig implementation achieves the fastest times across all configurations:

- **2.4× faster** than Rust for vacuum oscillations (25.6 ns vs 61 ns)
- **1.3× faster** than Rust for matter oscillations (73.3 ns vs 95 ns)
- **1.8× faster** than C++ for matter calculations

This performance advantage stems from:

1. **No ownership overhead:** Zig's explicit memory model allows aggressive inlining without borrow-checker constraints
2. **Direct SIMD:** Zig's `@Vector` type provides native SIMD operations
3. **Zero hidden allocations:** All computation occurs on the stack
4. **Simpler optimizer targets:** LLVM can optimize Zig's simpler IR more effectively

5.1.2. Rust Outperforms C++ for Matter Calculations

Key Result. Rust achieves **27% speedup** over C++ for matter oscillations with $N_{\text{Newton}} = 0$. This exceeds Dr. Denton's quoted 100 ns benchmark.

This advantage likely stems from:

1. **Ownership-enabled optimization:** Rust's strict aliasing rules allow LLVM to optimize more aggressively
2. **Loop vectorization:** Newton iteration inner loops optimize well in LLVM
3. **Zero-cost abstractions:** Rust idioms compile to efficient machine code

5.1.3. Vacuum Performance

For vacuum (simpler computation), C++ and Fortran are 20% faster. This is expected for compute kernels where Fortran excels.

5.1.4. Throughput

Language	Vacuum (M/s)	Matter (M/s)
Rust	17.5	10.5
C++	20.3	7.7
Fortran	19.7	9.4
Python	0.07	0.05

Table 4: Throughput in millions of calculations per second.

6. Applications

6.1. Interactive Visualization: Imagining the Neutrino

The `nufast-wasm` module powers *Imagining the Neutrino*, an interactive web-based educational tool:

<https://planckeon.github.io/itn/>

Features include:

- Real-time oscillation probability animation
- PMNS matrix visualization (2D and 3D representations)
- Energy spectrum and baseline scan plots
- CP asymmetry visualization: $A_{\text{CP}} = P(\nu_\mu \rightarrow \nu_e) - P(\bar{\nu}_\mu \rightarrow \bar{\nu}_e)$
- PREM Earth density model with automatic density calculation
- 11 experimental presets (DUNE, T2K, NOvA, etc.)
- Internationalization (7 languages)

With WASM, the visualization computes 400-point energy spectra in real-time (72 μs with VacuumBatch) as users adjust parameters.

6.2. Advanced Features (Zig v0.5.0)

The Zig implementation includes several advanced physics capabilities:

6.2.1. PREM Earth Model

Variable density calculations using the Preliminary Reference Earth Model (PREM):

```
// Automatic path integration through Earth layers
const probs = nufast.matterProbabilityPrem(params, 1300.0, 2.5);

// Get average density for a baseline
const avg = nufast.getAverageDensityAlongPath(1300.0);
```

The PREM implementation includes 6 layers (inner core, outer core, lower mantle, transition zone, upper mantle, crust) with appropriate densities and electron fractions. For a given baseline, the code calculates the path-weighted average density along the neutrino trajectory.

6.2.2. Non-Standard Interactions (NSI)

Support for new physics via NSI parameters:

```
const nsi = @import("nsi");
```

```

var nsi_params = nsi.NsiParams{
    .eps_ee = 0.1,      // Diagonal: enhanced electron potential
    .eps_em = nsi.Complex.init(0.03, 0.01), // Complex off-diagonal
};

const probs = nsi.matterProbabilityNsi(matter_nsi, 1300.0, 2.5);

```

NSI modifies the matter Hamiltonian: $H_{\text{matter}} \rightarrow A \times (\text{diag}(1, 0, 0) + \varepsilon)$ where ε is a Hermitian matrix. The implementation is accurate for $|\varepsilon| \leq 0.3$ (typical experimental bounds).

6.2.3. Sterile Neutrinos (3+1 Model)

Exact 4-flavor vacuum oscillations for short-baseline anomalies:

```

const sterile = @import("sterile");

const params = sterile.SterileParams.default;
const probs = sterile.sterileProbabilityVacuum(params, 500.0, 0.03);
// 4x4 probability matrix with active + sterile flavors

```

The sterile module implements the full 4×4 PMNS matrix with additional mixing angles ($\theta_{14}, \theta_{24}, \theta_{34}$) and CP phases (δ_{14}, δ_{24}). Note that the NuFast approximation does not extend to 4-flavor matter effects—the sterile module uses exact vacuum oscillation formulas.

6.3. Future Directions

- **GPU acceleration:** CUDA/WebGPU kernels for massive parallelism
- **4-flavor matter:** Exact numerical diagonalization for sterile + matter
- **ARM benchmarks:** Apple Silicon performance characterization

6.4. WebAssembly Deployment

The Zig implementation compiles to WebAssembly, enabling browser-based neutrino physics:

```

import { loadNuFast } from '@nufast/wasm';

const nufast = await loadNuFast();
nufast.setDefaultParams();

// Single-point calculation
const Pme = nufast.vacuumPmeDefault(1300, 2.5);

// Batch mode (2x faster)
const energies = new Float64Array(1000);
nufast.initVacuumBatch();
const results = nufast.vacuumBatchPme(1300, energies);

```

WASM performance (Bun runtime):

Mode	Single-point	Batch (1000)	Speedup
Vacuum	100 ns	50 ns/point	2×
Matter	150 ns	110 ns/point	1.4×

Table 5: WebAssembly performance. Batch mode amortizes JS↔WASM call overhead.

Key features:

- **Tiny binary:** 13.6 KB (baseline), 13.4 KB (SIMD128)
- **Zero dependencies:** No libc, pure math
- **TypeScript bindings:** Full type definitions with `NuFast` class
- **Batch API:** Pre-allocated 1024-point buffers for throughput

7. Conclusion

We have demonstrated that modern systems languages provide excellent platforms for computational neutrino physics:

Zig SIMD	48M vacuum/s, 27M matter/s (fastest)
Rust	27% faster than C++ for matter effects
WASM	13.6 KB binary, 20M vacuum/s batch mode
Memory safety	Compile-time guarantees (Rust), explicit control (Zig)
Distribution	crates.io (Rust), npm-ready (WASM)

Table 6: Summary of nufast capabilities.

The combination of performance, safety, and modern tooling—including WebAssembly for browser deployment—makes Rust and Zig attractive choices for future neutrino physics software development.

Acknowledgments

The author thanks Dr. Peter B. Denton (Brookhaven National Laboratory) for recommending the NuFast algorithm and providing guidance on neutrino oscillation calculation approaches. This work builds on the NuFast algorithm developed by Denton and Parke, with original implementations available at <https://github.com/PeterDenton/NuFast>.

The author acknowledges the neutrino physics education received during undergraduate research at Krea University under Dr. Sushant Raut, which provided the theoretical foundation for this implementation.

Code Availability

All code is open source under the MIT license:

Crates.io	https://crates.io/crates/nufast
GitHub	https://github.com/planckeon/nufast
Docs	https://docs.rs/nufast
Visualization	https://planckeon.github.io/itn/

Benchmark implementations (Rust, C++, Fortran, Python) are included in `benchmarks/`.

References

Bibliography

- [1] P. B. Denton, “Neutrino Oscillation Probabilities: A Compact Multi-algorithm Approach,” 2024.

- [2] P. B. Denton, S. J. Parke, and X. Zhang, “Eigenvector-based approach to neutrino oscillation probabilities in matter,” *Physical Review D*, vol. 101, p. 93001, 2020.
- [3] A. Abdullahi and S. J. Parke, “Eigenvalue-independent eigenvectors and the adjugate of a matrix,” 2022.
- [4] E. K. Akhmedov, R. Johansson, M. Lindner, T. Ohlsson, and T. Schwetz, “Series expansions for three-flavor neutrino oscillation probabilities in matter,” *JHEP*, vol. 4, p. 78, 2004.
- [5] C. A. Argüelles and B. J. P. Jones, “Neutrino Oscillations in a Quantum Computer,” *Physical Review D*, vol. 99, p. 96005, 2019.
- [6] J. Turro and others, “A Quantum simulation of neutrino-matter collective oscillations,” 2021.
- [7] P. B. Denton, H. Minakata, and S. J. Parke, “Compact Perturbative Expressions for Neutrino Oscillations in Matter.” 2016.
- [8] I. Esteban, M. C. González-García, M. Maltoni, T. Schwetz, and A. Zhou, “The fate of hints: updated global analysis of three-flavor neutrino oscillations.” 2020.